

Exercise 10

find the address of the test_backtrace function in obj/kern/kernel.asm, set a breakpoint there, and examine what happens each time it gets called after the kernel starts.

f0100040 <test_backtrace>:

The 2nd call to Test_backtrace ==> test_backtrace(4)

```
(gdb) p $esp
$3 = (void *) 0xf010ffb0
(gdb) p $ebp
$4 = (void *) 0xf010ffd8
```

The 3rd call to Test_backtrace ==> test_backtrace(3)

```
(gdb) p $esp
$6 = (void *) 0xf010ff9c
(gdb) p $ebp
$7 = (void *) 0xf010ffb8
```

The 4th call to Test_backtrace ==> test_backtrace(2)

```
(gdb) p $esp
$8 = (void *) 0xf010ff7c
(gdb) p $ebp
$9 = (void *) 0xf010ff98
```

The 5th call to Test_backtrace ==> test_backtrace(1)

```
(gdb) p $esp
$10 = (void *) 0xf010ff5c
(gdb) p $ebp
$11 = (void *) 0xf010ff78
```

The last call to Test_backtrace ==> test_backtrace(0)

```
(gdb) p $esp
$12 = (void *) 0xf010ff3c
(gdb) p $ebp
$13 = (void *) 0xf010ff58
```

How many 32-bit words does each recursive nesting level of test_backtrace push on the stack, and what are those words?

The difference in stack pointer between each successive calls of the function test_backtrace is 0x20 or 32 bytes.

It means each call of the test_backtrace makes the stack decreases by 32.

```
f0100040:      55                push    %ebp
f0100041:      89 e5            mov     %esp,%ebp
f0100043:      53              push    %ebx
f0100044:      83 ec 0c        sub     $0xc,%esp
f0100047:      8b 5d 08        mov     0x8(%ebp),%ebx
      cprintf("entering test_backtrace %d\n", x);
f010004a:      53              push    %ebx
f010004b:      68 20 1a 10 f0  push    $0xf0101a20
f0100050:      e8 f4 09 00 00  call    f0100a49 <cprintf>
      if (x > 0)
f0100055:      83 c4 10        add     $0x10,%esp
f0100058:      85 db          test    %ebx,%ebx
f010005a:      7e 11          jle     f010006d <test_backtrace+0x2d>
      test_backtrace(x-1);
f010005c:      83 ec 0c        sub     $0xc,%esp
f010005f:      8d 43 ff        lea     -0x1(%ebx),%eax
f0100062:      50              push    %eax
f0100063:      e8 d8 ff ff ff  call    f0100040 <test_backtrace>
```

Exercise 11

By studying kern/entry.S you'll find that there is an easy way to tell when to stop.

```
movl    $0x0,%ebp          # nu

# Set the stack pointer
movl    $(bootstacktop),%esp

# now to C code
call    i386_init
```

(Why can't the backtrace code detect how many arguments there actually are? How could this limitation be fixed?)

```
// or if there was no connecting function.
if (lfun < rfun)
    for (lline = lfun + 1;
         lline < rfun && stabs[lline].n_type == N_PSYM;
         lline++)
        info->eip_fn_narg++;
```

Then we used a switch statement to print the argument of a function.

```
switch(i.eip_fn_narg) {
    case 0:
        cprintf("EBP :%08x ,EIP %08x ,args: non \n",ptr,*(ptr+1));
        break;

    case 1:
        cprintf("EBP :%08x ,EIP %08x ,args: %08x \n",ptr,*(ptr+1),*(ptr+2));
        break;

    case 2:
        cprintf("EBP :%08x ,EIP %08x ,args: %08x , %08x \n",ptr,*(ptr+1),*(ptr+2),*(ptr+3));
        break;

    case 3:
        cprintf("EBP :%08x ,EIP %08x ,args: %08x , %08x , %08x \n",ptr,*(ptr+1),*(ptr+2),*(ptr+3), *(ptr+4));
        break;

    case 4:
        cprintf("EBP :%08x ,EIP %08x ,args: %08x , %08x , %08x , %08x \n",ptr,*(ptr+1),*(ptr+2),*(ptr+3), *(ptr+4), *(ptr+5));
        break;

    default: //5 or more
        cprintf("EBP :%08x ,EIP %08x ,args: %08x , %08x , %08x , %08x , %08x \n",ptr,*(ptr+1),*(ptr+2),*(ptr+3), *(ptr+4), *(ptr+5), *(ptr+6));
        break;
```

Implement the backtrace function as specified above.

```
EBP :f010ff38 ,EIP f010007b ,args: 00000000 , 00000001, f010ff78 , 00000000, f01008fd
EBP :f010ff58 ,EIP f0100068 ,args: 00000001 , 00000002, f010ff98 , 00000000, f01008fd
EBP :f010ff78 ,EIP f0100068 ,args: 00000002 , 00000003, f010ffb8 , 00000000, f01008fd
EBP :f010ff98 ,EIP f0100068 ,args: 00000003 , 00000004, 00000000 , 00000000, 00000000
EBP :f010ffb8 ,EIP f0100068 ,args: 00000004 , 00000005, 00000000 , 00010094, 00010094
EBP :f010ffd8 ,EIP f0100068 ,args: 00000005 , 00001aac, 00000644 , 00000000, 00000000
EBP :f010fff8 ,EIP f01000d4 ,args: 00111021 , 00000000, 00000000 , 00000000, 00000000
```

Exercise 12. Modify your stack backtrace function to display, for each eip, the function name, source file name, and line number corresponding to that eip.

```
EBP :f010ff38 ,EIP f010007b ,args: 00000000
Source File : kern/init.c Line# : 45 Func Name : test_backtrace:F(0,20) number of arguments : 1
EBP :f010ff58 ,EIP f0100068 ,args: 00000001
Source File : kern/init.c Line# : 28 Func Name : test_backtrace:F(0,20) number of arguments : 1
EBP :f010ff78 ,EIP f0100068 ,args: 00000002
Source File : kern/init.c Line# : 28 Func Name : test_backtrace:F(0,20) number of arguments : 1
EBP :f010ff98 ,EIP f0100068 ,args: 00000003
Source File : kern/init.c Line# : 28 Func Name : test_backtrace:F(0,20) number of arguments : 1
EBP :f010ffb8 ,EIP f0100068 ,args: 00000004
Source File : kern/init.c Line# : 28 Func Name : test_backtrace:F(0,20) number of arguments : 1
EBP :f010ffd8 ,EIP f0100068 ,args: 00000005
Source File : kern/init.c Line# : 28 Func Name : test_backtrace:F(0,20) number of arguments : 1
EBP :f010fff8 ,EIP f01000d4 ,args: non
Source File : kern/init.c Line# : 52 Func Name : i386_init:F(0,20) number of arguments : 0
```

In debuginfo_eip, where do __STAB_* come from?

- see if the bootloader loads the symbol table in memory as part of loading the kernel binary

```
moha@moha:~/6.828/lab$ objdump -h obj/kern/kernel
obj/kern/kernel: file format elf32-i386

Sections:
Idx Name          Size      VMA           LMA           File off  Algn
  0 .text          00001921 f0100000 00100000 00001000 2**4
  1 .rodata        000007d0 f0101940 00101940 00002940 2**5
  2 .stab          000039d9 f0102110 00102110 00003110 2**2
  3 .stabstr       000018fa f0105ae9 00105ae9 00006ae9 2**0
  4 .data          0000a300 f0108000 00108000 00009000 2**12
```

```
2 .stab          000039d9 f0102110 00102110 00003110 2**2
  CONTENTS, ALLOC, LOAD, READONLY, DATA
3 .stabstr       000018fa f0105ae9 00105ae9 00006ae9 2**0
  CONTENTS, ALLOC, LOAD, READONLY, DATA
4 .data          0000a300 f0108000 00108000 00009000 2**12
```

```
(gdb) x/20x 0x00102110
0x102110: 0x00000001 0x04d10000 0x000018f9 0x00000001
0x102120: 0x00000064 0xf0100000 0x00000012 0x00000084
0x102130: 0xf010000c 0x00000000 0x002c0044 0xf010000c
0x102140: 0x00000000 0x00390044 0xf0100015 0x00000000
0x102150: 0x003a0044 0xf010001a 0x00000000 0x003c0044
```

```
moha@moha:~/6.828/lab$ objdump -G obj/kern/kernel
obj/kern/kernel: file format elf32-i386

Contents of .stab section:

Symnum n_type n_othr n_desc n_value n_strx String
-----
-1      HdrSym 0      1233  000018f9 1
0       SO     0       0      f0100000 1      {standard input}
1       SOL    0       0      f010000c 18     kern/entry.S
2       SLINE  0       44     f010000c 0
3       SLINE  0       57     f0100015 0
4       SLINE  0       58     f010001a 0
```

Complete the implementation of `debuginfo_eip` by inserting the call to `stab_binsearch` to find the line number for an address.

```
stab_binsearch(stabs, &lline, &rline, N_SLINE, addr);
info->eip_line = stabs[lline].n_value;
```

Add a backtrace command to the kernel monitor.

```
static struct Command commands[] = {
    { "help", "Display this list of commands", mon_help },
    { "kerninfo", "Display information about the kernel", mon_kerninfo },
    { "backtrace", "Provides the backtrace", mon_backtrace},
};
```

```
K> help
help - Display this list of commands
kerninfo - Display information about the kernel
backtrace - Provides the backtrace
```

```
K> backtrace
EBP :f010ffd8 ,EIP f01009be ,args: 00000000
Source File : kern/monitor.c   Line# : 256   Func Name   : monitor:F(0,20)  number of arguments : 1

EBP :f010fff8 ,EIP f01000e1 ,args: non
Source File : kern/init.c     Line# : 67   Func Name   : i386_init:F(0,20)  number of arguments : 0
```

extend your implementation of `mon_backtrace` to call `debuginfo_eip` and print a line for each stack frame of the form:

```
K> backtrace
Stack backtrace:
  ebp f010ff78  eip f01008ae  args 00000001 f010ff8c 00000000 f0110580 00000000
    kern/monitor.c:143: monitor+106
```

```
address = *(ptr+1);

ptr1 = (uint32_t*) *ptr;
debuginfo_eip(address, &i);
```

```
EBP :f010ff38 ,EIP f010007b ,args: 00000000
Source File : kern/init.c   Line# : 45   Func Name   : test_backtrace:F(0,20)  number of arguments : 1

EBP :f010ff58 ,EIP f0100068 ,args: 00000001
Source File : kern/init.c   Line# : 28   Func Name   : test_backtrace:F(0,20)  number of arguments : 1

EBP :f010ff78 ,EIP f0100068 ,args: 00000002
Source File : kern/init.c   Line# : 28   Func Name   : test_backtrace:F(0,20)  number of arguments : 1

EBP :f010ff98 ,EIP f0100068 ,args: 00000003
Source File : kern/init.c   Line# : 28   Func Name   : test_backtrace:F(0,20)  number of arguments : 1

EBP :f010ffb8 ,EIP f0100068 ,args: 00000004
Source File : kern/init.c   Line# : 28   Func Name   : test_backtrace:F(0,20)  number of arguments : 1

EBP :f010ffd8 ,EIP f0100068 ,args: 00000005
Source File : kern/init.c   Line# : 28   Func Name   : test_backtrace:F(0,20)  number of arguments : 1

EBP :f010fff8 ,EIP f01000d4 ,args: non
Source File : kern/init.c   Line# : 52   Func Name   : i386_init:F(0,20)  number of arguments : 0
```


The rest of exercise 10

Before test called

```
esp    0xf010ffe0
ebp    0xf010fff8
```

After the call of backtrace , bcz the eip has been pushed onto the stack

```
esp    0xf010ffdc
ebp    0xf010fff8
```

To confirm that, we can read the memory of 0xf010ffdc, which is the address of the instruction after the call

```
(gdb) x/x 0xf010ffdc
0xf010ffdc:    0xf01000d4
```

```
f01000c8:    c7 04 24 05 00 00 00    movl    $0x5, (%esp)
f01000cf:    e8 6c ff ff ff         call    f0100040 <test_backtrace>
f01000d4:    83 c4 10              add     $0x10, %esp
```

```
f0100040 <test_backtrace>:
#include <kern/console.h>

// Test the stack backtrace function (lab 1 only)
void
test_backtrace(int x)
{
f0100040:    55                    push    %ebp
f0100041:    89 e5                mov     %esp, %ebp
f0100043:    53                    push    %ebx
f0100044:    83 ec 0c             sub     $0xc, %esp
f0100047:    8b 5d 08             mov     0x8(%ebp), %ebx
    cprintf("entering test_backtrace %d\n", x);
f010004a:    53                    push    %ebx
f010004b:    68 20 1a 10 f0       push    $0xf0101a20
f0100050:    e8 f4 09 00 00       call    f0100a49 <cprintf>
```

After the push instruction

%esp will be decreased by 4, so it becomes \Rightarrow 0xf010ffd8

And %ebp will remain the same \Rightarrow 0xf010fff8, but this value is now pushed on the stack as we can see by checking the memory of 0x...d8

```
(gdb) x/x 0xf010ffd8
0xf010ffd8:    0xf010fff8
```

After the mov instruction

After the mov instruction, %esp will remain the same 0xf010ffd8, but now %ebp will have the same value too.

%ebx 65684

After push %ebx

%esp decreases by 4 and becomes 0xf010ffd4,

Because Ebx is a Callee saved register (EBX, ESI & EDI), and since we are going to use it, then we have to save it.

sub \$0xc,%esp

Esp becomes 0xf010ffc8, this instructions it will give some storage to the local variables and complicated operations.

Keep in mind that the Test_Backtrace function is written in C, so these instructions are done by the compiler.

mov 0x8(%ebp),%ebx

By convention, the first argument is always stored here, so this will put x which is 5 into ebx

Note that x, was pushed by the function that called Test_backtrace

```
f01000c8:  c7 04 24 05 00 00 00    movl    $0x5, (%esp)
f01000cf:  e8 6c ff ff ff          call    f0100040 <test_backtrace>
f01000d4:  83 c4 10                add     $0x10, %esp
```

Now we will start preparing to call the function printf(), so we need first to save its 2 arguments (x and the string on the stack)

```
f010004a:  53                      push    %ebx
f010004b:  68 80 18 10 f0          push    $0xf0101880
f0100050:  e8 95 08 00 00          call    f01008ea <printf>
```

So the esp will decrease by 12, and becomes 0xf010ffc8, %ebp still the same,

But let's check the contents of the stack:

It shd be the %ebp address f0100055, then the second argument 0xf...880 , then the first argument x which is 5, then 4 garbage words bcz of the subtraction.

```
0xf010ffbc:  0xf0100055  0xf0101880  0x00000005  0x00000000
0xf010ffcc:  0x00010094  0x00010094  0x00010094  0xf010fff8
0xf010ffdc:  0xf01000d4  0x00000005
```

Now the Printf Prologue, which is saving the base pointer of the Test_Backtrace, then assigning a new base pointer to the current function.

Both `%esp` and `%ebp` will be `0xf010ffb4`, and the value saved at this stack address is `0xf010ffd8` which is the base pointer of the `Test_Backtrace`.
Then subtract 16 from the `esp`, so it becomes `0xf010ffa4`

```

K> kerninfo
Special kernel symbols:
_start      0010000c (phys)
entry      f010000c (virt) 0010000c (phys)
etext      f0101941 (virt) 00101941 (phys)
edata      f0112300 (virt) 00112300 (phys)
end        f0112944 (virt) 00112944 (phys)
Kernel executable memory footprint: 75KB

```

Symnum	n_type	n_othr	n_desc	n_value	n_strx	String
-1	HdrSym	0	1233	000018f9	1	
0	SO	0	0	f0100000	1	{standard input}
1	SOL	0	0	f010000c	18	kern/entry.S
2	SLINE	0	44	f010000c	0	
3	SLINE	0	57	f0100015	0	
4	SLINE	0	58	f010001a	0	
5	SLINE	0	60	f010001d	0	
6	SLINE	0	61	f0100020	0	
7	SLINE	0	62	f0100025	0	
8	SLINE	0	67	f0100028	0	
9	SLINE	0	68	f010002d	0	
10	SLINE	0	74	f010002f	0	
11	SLINE	0	77	f0100034	0	
12	SLINE	0	80	f0100039	0	
13	SLINE	0	83	f010003e	0	
14	SO	0	2	f0100040	31	kern/entrypgdir.c
15	OPT	0	0	00000000	49	gcc2_compiled.
16	LCVM	0	0	00000000	64	input(0..1) = 0..1; 341740

```

stab begin = f0102150
stab end = f0105b40

```

```

stab begin = f0102230
stab end = f0105c68
stabstr begin = f0105c69
stabstr end = f0107562

```

308	FUN	0	0	f01004e9	3193	kbd_intr:F(0,20)
309	SLINE	0	365	00000000	0	
310	SLINE	0	366	00000006	0	
311	SLINE	0	367	00000010	0	
312	FUN	0	0	f01004fb	3210	cons_getc:F(0,1)
313	SLINE	0	408	00000000	0	
314	SLINE	0	414	00000006	0	

n_value for a function contains its address

But for an SLINE, it contains its offset i believe

```
rfile = 4d9
lfunc = 63
rfunc = 6c
12addr = 28
```

EBP :f010ff18	,EIP f010007b	,args: 00000000 , 00000000
EBP :f010ff38	,EIP f0100068	,args: 00000000 , 00000001
EBP :f010ff58	,EIP f0100068	,args: 00000001 , 00000002
EBP :f010ff78	,EIP f0100068	,args: 00000002 , 00000003
EBP :f010ff98	,EIP f0100068	,args: 00000003 , 00000004
EBP :f010ffb8	,EIP f0100068	,args: 00000004 , 00000005
EBP :f010ffd8	,EIP f01000d4	,args: 00000005 , 00001aac

What are these 32?

Test_backtrace ()

Push both ebx and ebp besides subtracting the stack by 12 ⇒ these are 16 bytes

Cprintf()

Pushing the ebx and the string ⇒ 8 bytes

```
int
cprintf(const char *fmt, ...)
{
    va_list ap;
    int cnt;

    va_start(ap, fmt);
    cnt = vcprintf(fmt, ap);
    va_end(ap);

    return cnt;
}
```

```
int
cprintf(const char *fmt, ...)
{
f0100a49:      55                |      push    %ebp
f0100a4a:      89 e5            |      mov     %esp,%ebp
f0100a4c:      83 ec 10         |      sub     $0x10,%esp
    va_list ap;
    int cnt;

    va_start(ap, fmt);
f0100a4f:      8d 45 0c         |      lea     0xc(%ebp),%eax
    cnt = vcprintf(fmt, ap);
f0100a52:      50              |      push    %eax
f0100a53:      ff 75 08         |      pushl   0x8(%ebp)
f0100a56:      e8 c8 ff ff ff   |      call    f0100a23 <vcprintf>
    va_end(ap);

    return cnt;
}
```

The following is the preparing to call vcprintf(), so its 2 arguments are pushed on the stack (

```

        cnt = vcprintf(fmt, ap);
f01008f3:      50                push    %eax
f01008f4:      ff 75 08        pushl   0x8(%ebp)
f01008f7:      e8 c8 ff ff ff    call    f01008c4 <vcprintf>

```

va_list ap; is the first argument so it is at 8+ebp, the 2nd argument fmt is saved by the previous instruction in eax. Then call the function vcprintf
 %esp decreased by 12, so will be 0xf010fef8 and %ebp will still be 0xf010ffb4

```

(gdb) p /x $eax
$11 = 0xf010ffc4

```

```

0xf010ff9c:  0xf01008fc  0xf0101880  0xf010ffc4  0x00000000
0xf010ffac:  0x00000000  0x00000000  0x00000000  0xf010ffd8
0xf010ffbc:  0xf0100055  0xf0101880

```

```

(gdb) x/30x $esp
0xf010ff9c:  0xf01008fc  0xf0101880  0xf010ffc4  0x00000000
0xf010ffac:  0x00000000  0x00000000  0x00000000  0xf010ffd8
0xf010ffbc:  0xf0100055  0xf0101880  0x00000005  0x00000000
0xf010ffcc:  0x00010094  0x00010094  0x00010094  0xf010fff8
0xf010ffdc:  0xf01000d4  0x00000005  0x00001aac  0x00000644
0xf010ffec:  0x00000000  0x00000000  0x00000000  0x00000000
0xf010fffc:  0xf010003e  0x00111021  0x00000000  0x00000000
0xf011000c <entry_pgdir+12>:  0x00000000  0x00000000

```

after making the new frame of vcprintf()

```

(gdb) p $esp
$12 = (void *) 0xf010ff80
(gdb) p $ebp
$13 = (void *) 0xf010ff98

```

Then it calls vprintfmt with 4 arguments, so 4 words are pushed on the stack

Then vmprintf stores the 3 callee registers

Then it pushes ebx and eax in there...

```

f0100ccd:  53                push    %ebx
f0100cce:  50                push    %eax
f0100ccf:  ff d6            call    *%esi

```

How many 32-bit words does each recursive nesting level of test_backtrace push on the stack, and **what are those words**?

```
mov %ebp, esp
```

```
pop %ebp
```

```
CALL, (push eip)
```

```
Ret instruction => pop %Eip
```

```
f01008ea <printf>
```

```
f0100040 <test_backtrace>:
```